

Automation

Ansible Conditionals and Loops



**DE HOGESCHOOL
MET HET NETWERK**

Elfde-Liniestraat 24, 3500 Hasselt, www.pxl.be

Conditionals

Types of Ansible Conditionals

- **When Conditionals:** most common conditionals in Ansible - they specify conditions that must be met for a task to be executed.
- **Failed Conditionals:** specify tasks that should be executed if a previous task fails.
- **Changed Conditionals:** specify tasks that should be executed if a previous task has made changes to the system.
- **Custom Conditionals:** conditionals that you define yourself using Ansible's Jinja2 template language.

```
- name: Do something if myvar is set  
  command: /path/to/some/command  
  when: myvar == "somevalue"
```

Understanding Ansible Conditionals

- Ansible conditionals are used to make decisions in playbooks based on certain conditions.
- For example, you might want to run a task only if a certain variable is set to a specific value.
- Conditionals are expressed using the **when** keyword.

```
- name: Install Apache if it's not already installed
  ansible.builtin.package:
    name: apache2
    state: present
  when: "'apache2' not in ansible_facts.packages"
```

Ansible Conditionals operators

- Comparison operators (e.g., ==, >, <, !=)
- Logical operators (e.g., **and**, **or**, **not**)
- Regular expression matching (e.g., =~)

```
tasks:
```

```
- name: Shut down CentOS 6 and Debian 7 systems
```

```
  ansible.builtin.command: /sbin/shutdown -t now
```

```
  when: (ansible_facts['distribution'] == "CentOS" and ansible_facts['distribution_major_version'] == "6") or  
        (ansible_facts['distribution'] == "Debian" and ansible_facts['distribution_major_version'] == "7")
```

Implicit logical operator **and**

Specify multiple conditions that all need to be true as a list (== a logical and)

```
tasks:
```

```
- name: Shut down CentOS 6 systems
```

```
  ansible.builtin.command: /sbin/shutdown -t now
```

```
  when:
```

```
    - ansible_facts['distribution'] == "CentOS"
```

```
    - ansible_facts['distribution_major_version'] == "6"
```

Testing dictionary membership with **in**

example playbook to install vsftpd (very secure ftp daemon) on ftp servers.

```
---
- name: Install vsftpd on ftpservers
  hosts: all
  become: yes
  tasks:
    - name: Install FTP server when host in ftpserver group
      yum:
        name: vsftpd
        state: latest
      when: inventory_hostname in groups["ftpserver"]
```


Jinja2 templates and filters

- Filters let you transform JSON data into YAML data, split a URL to extract the hostname, get the SHA1 hash of a string, add or multiply integers, etc.
- List of built-in [filters](#) in the official Jinja2 template documentation.
- Providing default values
 - `{{ some_variable | default(5) }}`
 - `{{ lookup('env', 'MY_USER') | default('admin', true) }}`
- Casting data types
 - `when: some_string_value | bool`
 - `when: ansible_facts['lsb']['major_release'] | int >= 6`

Using registered variables

A registered variable is global, registered with the keyword **register**, and contains the status and the output of the task that created it.

Use **variable.stdout** to access the string contents of the variables.

```
tasks:
```

- name: Register a variable
ansible.builtin.shell: cat /etc/motd
register: motd_contents
- name: Use the variable in conditional statement
ansible.builtin.shell: echo "motd contains the word hi"
when: motd_contents.stdout.find('hi') != -1

Using registered variables

```
---
- name: Ansible conditional with a registered variable example
  hosts: web
  remote_user: ubuntu
  become: true
  tasks:
# to list the directory content in '/etc/hosts directory'
  - name: List contents of directory and Store in content1
    ansible.builtin.command: ls /etc/hosts
    register: contents1
# to list the directory content in '/home/ubuntu/hello'
  - name: List contents of directory and Store in content2
    ansible.builtin.command: ls /home/ubuntu/hello
    register: contents2
# display Directory is empty if the directories
# '/etc/hosts' or '/home/ubuntu/hello' are empty
  - name: Check contents for emptiness for content1 or content2
    ansible.builtin.debug:
      msg: "Directory is empty"
    when: contents1.stdout == "" or contents2.stdout == ""
```

Return values of registered variables

- Ansible modules return a data structure that can be registered into a variable, or seen directly when output by the ansible program.
- Each module can optionally document its own unique return values.
- https://docs.ansible.com/ansible/latest/reference_appendices/common_return_values.html

```
TASK [Second Task - Print the full output] *****
ok: [ubuntu.anslab.com] => {
  "virtualenv_output": {
    "ansible_facts": {
      "discovered_interpreter_python": "/usr/bin/python3"
    },
    "changed": true,
    "cmd": "which virtualenv",
    "delta": "0:00:00.003394",
    "end": "2022-09-21 21:08:57.717425",
    "failed": true,
    "msg": "non-zero return code",
    "rc": 1,
    "start": "2022-09-21 21:08:57.714031",
    "stderr": "",
    "stderr_lines": [],
    "stdout": "",
    "stdout_lines": []
  }
}
```

Use the status of registered variables: **failed, succeeded, skipped, changed**

```
tasks:
- name: Register a variable, ignore errors and continue
  ansible.builtin.command: /bin/false
  register: result
  ignore_errors: true

- name: Run only if the task that registered the "result" variable fails
  ansible.builtin.command: /bin/something
  when: result is failed

- name: Run only if the task that registered the "result" variable succeeds
  ansible.builtin.command: /bin/something_else
  when: result is succeeded

- name: Run only if the task that registered the "result" variable is skipped
  ansible.builtin.command: /bin/still/something_else
  when: result is skipped

- name: Run only if the task that registered the "result" variable changed something.
  ansible.builtin.command: /bin/still/something_else
  when: result is changed
```

Example of failed, succeeded

```
---
- hosts: all
  vars:
    - user: angie
  tasks:
    - name: Check if file already exists
      command: ls /home/{{ user }}/myfile
      register: file_exists
      ignore_errors: true

    - name: create file for user
      file:
        path: /home/{{ user }}/myfile
        state: touch
        when: file_exists is failed

    - name: show message if file exists
      debug:
        msg: The user file already exists.
        when: file_exists is succeeded
```

- this playbook checks if the file `/home/angie/myfile` exists on the target host.
- If the file exists, it displays a message.
- If the file does not exist, it creates the file.
- The **register** directive is used to capture the output of the command module, and the **when** directive is used to conditionally execute the file and debug modules based on the value of the `file_exists` variable.

defined and undefined

tasks:

- name: Run the command if "foo" is defined
 ansible.builtin.shell: echo "I've got '{{ foo }}' and am not afraid to use it!"
 when: foo is defined
- name: Fail if "bar" is undefined
 ansible.builtin.fail: msg="Bailing out. This play requires 'bar'"
 when: bar is undefined

failed_when and changed_when

- Ansible **failed_when** and **changed_when** statements are similar to ansible **when** statement. The only difference is that It will mark the task as failed or success[changed], when the condition defined, is met or satisfied.
- The primary purpose of the **failed_when** and **changed_when** statements are to determine whether the task is actually successful or results in a failure.

Defining failure

- The `failed_when` conditional lets you define what "failure" means in each task.
- As with all conditionals in Ansible, lists of multiple `failed_when` conditions are joined with an implicit **and**, meaning the task only fails when all conditions are met.
- If you want to trigger a failure when any of the conditions is met, define the conditions in a string with an explicit **or** operator.

```
- name: Fail task when the command error output prints Boom!  
  ansible.builtin.command: /usr/bin/example-command -x -y -z  
  register: command_result  
  failed_when: "'Boom !' in command_result.stderr"
```

```
- name: Check if a file exists in temp and fail task if it does  
  ansible.builtin.command: ls /tmp/this_should_not_be_here  
  register: result  
  failed_when:  
    - result.rc == 0  
    - '"No such" not in result.stdout'
```

```
- name: Fail task when both files are identical  
  ansible.builtin.raw: diff foo/file1 bar/file2  
  register: diff_cmd  
  failed_when: diff_cmd.rc == 0 or diff_cmd.rc >= 2
```

Tip: splitting conditions over multiple lines with >

If you have too many conditions to fit into one line, split it into a multi-line YAML value with >.

```
- name: example of many failed_when conditions with 'or'
  ansible.builtin.shell: "./mySuperProgram"
  register: my_return
  failed_when: >
    ("No such file or directory" in ret.stdout) or
    (ret.stderr != '') or
    (ret.rc == 10)
```

Ignoring failure with **ignore_errors**

- By default Ansible stops executing tasks on a host when a task fails on that host. You can use **ignore_errors** to continue on in spite of the failure.
- Only works when the task is able to run and returns a value of 'failed'.
- Does not make Ansible ignore undefined variable errors, connection failures, execution issues (for example, missing packages), or syntax errors.

```
- name: List a non-existent file
  command: ls non-existent.txt
  ignore_errors: true

- name: continue task after failing ls
  debug:
    msg: "Continue task after failure"
```

Ignoring unreachable host errors with **ignore_unreachable**

- If Ansible cannot connect to a host, it marks that host as 'UNREACHABLE' and removes it from the list of active hosts for the run.
- You can ignore a task failure due to the host instance being 'UNREACHABLE' with the **ignore_unreachable** keyword.
- Can be defined on playbook as well as on task level.
- With **ignore_unreachable** Ansible ignores the task errors, but continues to execute future tasks against the unreachable host.

```
- name: This executes, fails, and the failure is ignored
  ansible.builtin.command: /bin/true
  ignore_unreachable: true

- name: This executes, fails, and ends the play for this host
  ansible.builtin.command: /bin/true
```

Defining changed

- The **changed_when** conditional lets you define when a particular task has "changed" a remote node.
- This lets you determine, based on return codes or output, whether a change should be reported in Ansible statistics and whether a **handler** (see *later*) should be triggered or not.

```
- name: Report 'changed' when the return code is not equal to 2
  ansible.builtin.shell: /usr/bin/mybinary --mode="analyze"
  register: binary_result
  changed_when: "binary_result.rc != 2"

- name: This will never report 'changed' status
  ansible.builtin.shell: wall 'beep'
  changed_when: False
```

Selecting variables files based on `ansible_facts`

```
---
- hosts: webservers
  remote_user: root
  vars_files:
    - "vars/common.yml"
    - [ "vars/{{ ansible_facts['os_family'] }}.yml", "vars/os_defaults.yml" ]
  tasks:
    - name: Make sure apache is started
      ansible.builtin.service:
        name: '{{ apache }}'
        state: started
```

```
---
# for vars/RedHat.yml
apache: httpd
somethingelse: 42
```

```
---
# for vars/Debian.yml
apache: apache2
somethingelse: 13
```

Grouping tasks with **block**:

- **block** in an Ansible playbook is used to group multiple tasks together.
- useful for applying the same attributes or error handling to a set of tasks, making the playbook easier to read and maintain.
- E.g., use a **block** to group tasks that should be executed together under a specific condition. You can then apply the **when** keyword to the entire block, rather than repeating the condition for each individual task.
- In an Ansible playbook, **rescue** is used to specify tasks that should be executed when any task within a **block** fails, while **always** defines tasks that will be executed after the block and rescue sections, regardless of their success or failure

```
tasks:
  - name: Deploy and configure Nginx
    block:
      - name: Copy Nginx configuration
        ansible.builtin.template:
          src: nginx.conf.j2
          dest: "{{ nginx_config_file }}"
          owner: root
          group: root
          mode: 0644

      - name: Enable site configuration
        ansible.builtin.file:
          src: "{{ nginx_config_file }}"
          dest: /etc/nginx/sites-enabled/{{ app_name }}
          state: link

      - name: Reload Nginx
        ansible.builtin.systemd:
          name: nginx
          state: reloaded
    when: "'web_servers' in group_names"

  rescue:
    - name: Remove faulty configuration
      ansible.builtin.file:
        path: "{{ nginx_config_file }}"
        state: absent

    - name: Rollback Nginx to default site
      ansible.builtin.command: ln -sf /etc/nginx/sites-available/default /etc/nginx/sites-enabled/default
      args:
        removes: /etc/nginx/sites-enabled/{{ app_name }}

    - name: Restart Nginx
      ansible.builtin.systemd:
        name: nginx
        state: restarted

    - name: Fail the playbook
      ansible.builtin.fail:
        msg: "Nginx configuration failed. Rolled back to default site."

  always:
    - name: Check Nginx status
      ansible.builtin.systemd:
        name: nginx
        state: started
```


Loops

Ansible Loops

- Implement loops using **loop**, **with_<lookup>** (**with_list**, **with_items**, **with_dict**, ...), and **until**.
- Using **loop** is the recommended way vs **with_<lookup>**.
- Access each element of the looped-over list with "**{{ item }}**"

```
- name: Add several users
  ansible.builtin.user:
    name: "{{ item }}"
    state: present
    groups: "wheel"
  loop:
    - testuser1
    - testuser2
```

Looping over hashes

```
---
- name: Ensure users
  hosts: node1
  become: true

  tasks:
    - name: Ensure three users are present
      user:
        name: "{{ item.username }}"
        state: present
        groups: "{{ item.groups }}"
      loop:
        - { username: 'dev_user', groups: 'ftp' }
        - { username: 'qa_user', groups: 'ftp' }
        - { username: 'prod_user', groups: 'apache' }
```

Loop example

```
- name: Registered variable usage as a loop list
hosts: all
tasks:

  - name: Retrieve the list of home directories
    ansible.builtin.command: ls /home
    register: home_dirs

  - name: Add home dirs to the backup spooler
    ansible.builtin.file:
      path: /mnt/bkspool/{{ item }}
      src: /home/{{ item }}
      state: link
    loop: "{{ home_dirs.stdout_lines }}"
    # same as loop: "{{ home_dirs.stdout.split() }}"
```

- **home_dirs** is not a list, so difficult for looping.
- However, the Jinja method **.stdout_lines** returns a list of lines. Now we can use the list for iterating over with **loop**.

Looping dictionaries

```
users:
  - name: john
    uid: 1001
    gid: 1001
    password: mypassword
  - name: jane
    uid: 1002
    gid: 1002
    password: myotherpassword
```

```
- name: Import users file
hosts: myserver
vars_files:
  - /path/to/group_vars/prod/users.yml
tasks:
  - name: Create users with properties
    user:
      name: "{{ item.name }}"
      uid: "{{ item.uid }}"
      group: "{{ item.gid }}"
      password: "{{ item.password | password_hash('sha512') }}"
      state: present
    loop: "{{ users }}"
```

Combining Loops and Conditionals

- If you combine a when statement with a loop, Ansible processes the condition separately for each item.
- You can execute the task on some items in the loop and skip it on other items.
- You can just use the keyword **item** in the **when:** clause.

```
tasks:  
  - name: run with items greater than 5  
    ansible.builtin.command: echo {{ item }}  
    loop: [ 0, 2, 4, 6, 8, 10 ]  
    when: item > 5
```

Customize loops with `loop_control`:

- `loop_control` is used to modify how the loop executes and provides more control over the loop iteration.
- `loop_control` is a dictionary containing the following keys:
 - **label**: allows you to customize the label used to identify the item being processed in the loop.
 - **index_var**: allows you to specify a variable to hold the index of the current iteration.
 - **loop_var**: allows you to specify a variable to hold the value of the current iteration.
 - **extended**: allows you to specify additional loop options like **skip_missing**, **stop_execution**, etc.
 - **flatten**: allows you to flatten the items in the loop before iterating over them.
 - **pause**: allows you to pause the loop at a specified iteration and resume it later.

```
- name: Loop example
hosts: localhost
tasks:
  - name: Print numbers with custom label
    debug:
      msg: "The number is {{ item }}"
    loop:
      - 1
      - 2
      - 3
    loop_control:
      label: "Number {{ item }}"
```

```
TASK [Print numbers with custom label] *****
ok: [localhost] => (label=Number 1) => {
    "msg": "The number is 1"
}
ok: [localhost] => (label=Number 2) => {
    "msg": "The number is 2"
}
ok: [localhost] => (label=Number 3) => {
    "msg": "The number is 3"
}
```


Error handling with **until**, **retries** and **delay**

- **until**: a condition that must be met for a task or block of tasks to succeed. If the condition is not met, the task will retry until the condition is met or the maximum number of retries is exceeded.
- **retries**: maximum number of times to retry a task or block of tasks if they fail.
- **delay**: amount of time to wait between retries.

```
---
- name: Retry until a file is available
  hosts: localhost
  tasks:
    - name: Validate if the file is present
      shell: ls -lrt /tmp/myprocess.pid
      register: lsresult
      until: "lsresult is not failed"
      retries: 10
      delay: 10
```

```
PLAY [Retry until a file is available] *****

TASK [Gathering Facts] *****
ok: [localhost]

TASK [Validate if the file is present] *****
FAILED - RETRYING: Validate if the file is present (10 retries left).
FAILED - RETRYING: Validate if the file is present (9 retries left).
FAILED - RETRYING: Validate if the file is present (8 retries left).
FAILED - RETRYING: Validate if the file is present (7 retries left).
FAILED - RETRYING: Validate if the file is present (6 retries left).
FAILED - RETRYING: Validate if the file is present (5 retries left).
FAILED - RETRYING: Validate if the file is present (4 retries left).
FAILED - RETRYING: Validate if the file is present (3 retries left).
FAILED - RETRYING: Validate if the file is present (2 retries left).
FAILED - RETRYING: Validate if the file is present (1 retries left).
fatal: [localhost]: FAILED! => {"attempts": 10, "changed": true, "cmd": "ls -lrt /tmp/myprocess.pid", "delta": "0:00:00.007821", "end": "2022-04-07 02:46:49.676415", "msg": "non-zero return code", "rc": 1, "start": "2022-04-07 02:46:49.668594", "stderr": "ls: /tmp/myprocess.pid: No such file or directory", "stdout_lines": ["ls: /tmp/myprocess.pid: No such file or directory"], "stdout": "", "stdout_lines": []}

PLAY RECAP *****
localhost                : ok=1    changed=0    unreachable=0    failed=1    skipped=0
rescued=0                ignored=0
```

Exercises

Web server

Configure a web server on an ubuntu machine using Ansible.

The web server should have the following features:

- Apache web server installed and running
- PHP installed and configured to work with Apache and mysql
- ubuntu packages: apache2, php, libapache2-mod-php, php-mysql, mysql-server, python3-pip
- try 2 more times if package update fails
- A custom index.php file "<?php phpinfo(); ?>" created in the Apache document root
- The web server should be configured to listen on port 8080 instead of the default port 80
- A MySQL server installed and configured with a database and user
 - install pip file pymysql first. then use the mysql_user module.
- The firewall UFW should be configured to allow incoming traffic on port 8080 and port 3306. Use a loop.

end

~~node.js and postgres~~ too hard/messy

Create an Ansible playbook named `playbook.yml`. In this playbook, you will:

- Update the package cache.
- Install Node.js and npm on the Node.js application rhel 9 server.
- Install PostgreSQL on the PostgreSQL database rhel 9 server.
- Deploy the sample Node.js application `app.js` on the application server.
- Configure the PostgreSQL server with a new database and user for the Node.js application.
- Start and enable the Node.js application and the PostgreSQL service.

Use variables, conditionals, and loops in the Ansible playbook to accomplish the tasks for each server.

```
const http = require('http');
const hostname = '0.0.0.0';
const port = 3000;

const server = http.createServer((req, res) => {
  res.statusCode = 200;
  res.setHeader('Content-Type', 'text/plain');
  res.end('Hello, World!\n');
});

server.listen(port, hostname, () => {
  console.log(`Server running at
http://${hostname}:${port}/`);
});
```

end